

# Federated Machine Learning: Concept and Applications

QIANG YANG, Hong Kong University of Science and Technology, Hong Kong

YANG LIU and TIANJIAN CHEN, Webank, China

YONGXIN TONG, Beihang University, China

---

Today's artificial intelligence still faces two major challenges. One is that, in most industries, data exists in the form of isolated islands. The other is the strengthening of data privacy and security. We propose a possible solution to these challenges: secure federated learning. Beyond the federated-learning framework first proposed by Google in 2016, we introduce a comprehensive secure federated-learning framework, which includes horizontal federated learning, vertical federated learning, and federated transfer learning. We provide definitions, architectures, and applications for the federated-learning framework, and provide a comprehensive survey of existing works on this subject. In addition, we propose building data networks among organizations based on federated mechanisms as an effective solution to allowing knowledge to be shared without compromising user privacy.

**CCS Concepts:** • Security and privacy; • Computing methodologies → Artificial intelligence; Machine learning; Supervised learning;

**Additional Key Words and Phrases:** Federated learning, GDPR, transfer learning

**ACM Reference format:**

Qiang Yang, Yang Liu, Tianjian Chen, and Yongxin Tong. 2019. Federated Machine Learning: Concept and Applications. *ACM Trans. Intell. Syst. Technol.* 10, 2, Article 12 (January 2019), 19 pages.

<https://doi.org/10.1145/3298981>

---

## 1 INTRODUCTION

2016 is the year when artificial intelligence (AI) came of age. With AlphaGo [59] defeating the top human Go players, we have truly witnessed the huge potential in AI and have began to expect more complex, cutting-edge AI technology in many applications, including driverless cars, medical care, and finance. Today, AI technology is showing its strengths in almost every industry and most walks of life. However, when we look back at the development of AI, it is inevitable that it has experienced several ups and downs. Will there be a next downturn for AI? When will it appear and because of what factors? The current public interest in AI is partly driven by Big Data availability: AlphaGo in 2016 used a total of 300,000 games as training data to achieve excellent results.

With AlphaGo's success, people naturally hope that the big data-driven AI such as AlphaGo will be realized soon in all aspects of our lives. However, real-world situations are somewhat disappointing: with the exception of a few industries, most fields have only limited data or poor-quality

---

Authors' addresses: Q. Yang, Hong Kong University of Science and Technology, Clear Water Bay, Kowloon, Hong Kong, China; email: qyang@cse.ust.hk; Y. Liu and T. Chen, Building #7 A, Keji Shengtaiyuan, No. 1819, Shahe West Road, Nanshan District, Shenzhen, China; emails: {yangliu, tobychen}@webank.com; Y. Tong (corresponding author), Advanced Innovation Center for Big Data and Brain Computing, Beihang University, No. 37 Xueyuan Road, Haidian District, Beijing, China; email: yxtong@buaa.edu.cn.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

© 2019 Copyright held by the owner/author(s).

2157-6904/2019/01-ART12

<https://doi.org/10.1145/3298981>



Fig. 1. GDPR: EU regulation on data protection.

data, making the application of AI technology more difficult than we thought. Would it be possible to fuse the data together in a common site by transporting the data across organizations? In fact, it is very difficult, if not impossible, in many situations to break the barriers between data sources. In general, the data required in any AI project involves multiple types. For example, in an AI-driven product recommendation service, the product seller has information about the product and data of the user's purchase but not the data that describes the user's purchasing ability and payment habits. In most industries, data exists in the form of isolated islands. Due to industry competition, privacy security, and complicated administrative procedures, even data integration between different departments of the same company faces heavy resistance. It is almost impossible to integrate the data scattered around the country and institutions, or the cost is prohibited.

At the same time, with the increasing awareness of large companies compromising on data security and user privacy, the emphasis on data privacy and security has become a worldwide major issue. News about leaks on public data are causing great concern in public media and governments. For example, the recent data breach by Facebook has caused a wide range of protests [70]. In response, states across the world are strengthening laws to protect data security and privacy. An example is the General Data Protection Regulation (GDPR) [19] enforced by the European Union on May 25, 2018. The GDPR (Figure 1) aims to protect users' personal privacy and provide data security. It requires businesses to use clear and plain language for their user agreement and grants users the "right to be forgotten," that is, users can have their personal data deleted or withdrawn. Companies violating the bill will face stiff fines. Similar acts of privacy and security are being enacted in the United States and China. For example, China's Cyber Security Law and the General Principles of the Civil Law, enacted in 2017, require that Internet businesses must not leak or tamper with the personal information that they collect and that, when conducting data transactions with third parties, they need to ensure that proposed contracts follow legal data protection obligations. The establishment of these regulations will clearly help build a more civil society but will also pose new challenges to the data-transaction procedures commonly used today in AI.

To be more specific, traditional data-processing models in AI often involve simple data-transaction models, with one party collecting and transferring data to another party and the other party being responsible for cleaning and fusing the data. Finally a third party will take the integrated data and build models for still other parties to use. The models are usually the final products that are sold as a service. This traditional procedure face challenges with the above new data regulations and laws. As well, since users may be unclear about the future uses of the models, the transactions violate laws such as the GDPR. As a result, we face a dilemma that our data is in the form of isolated islands, but we are forbidden in many situations to collect, fuse, and use the data in different places for AI processing. How to legally solve the problem of data fragmentation and isolation is a major challenge for AI researchers and practitioners today.

In this article, we give an overview of a new approach, known as *federated learning*, which is a possible solution for these challenges. We survey existing works on federated learning, and propose definitions, categorizations, and applications for a comprehensive secure federated-learning framework. We discuss how the federated-learning framework can be applied to various businesses successfully. In promoting federated learning, we hope to shift the focus of AI development from improving model performance, which is what most of the AI field is currently doing, to investigating methods for data integration that are compliant with data privacy and security laws.

## 2 AN OVERVIEW OF FEDERATED LEARNING

The concept of federated learning was proposed by Google recently [36, 37, 41]. Google's main idea is to build machine-learning models based on datasets that are distributed across multiple devices while preventing data leakage. Recent improvements have been focusing on overcoming the statistical challenges [60, 77] and improving security [9, 23] in federated learning. There are also research efforts to make federated learning more personalizable [13, 60]. The above works all focus on on-device federated learning in which distributed mobile-user interactions are involved and communication cost in massive distribution, unbalanced data distribution, and device reliability are some of the major factors for optimization. In addition, data is partitioned by user IDs or device IDs, therefore, *horizontally* in the data space. This line of work is highly related to privacy-preserving machine learning, as reported in [58] because it also considers data privacy in a decentralized collaborative-learning setting. To extend the concept of federated learning to cover collaborative-learning scenarios among organizations, we extend the original "federated learning" to a general concept for all privacy-preserving decentralized collaborative machine-learning techniques. In [71], we have given a preliminary overview of the federated-learning and federated transfer-learning technique. In this article, we further survey the relevant security foundations and explore the relationship with several other related areas, such as multiagent theory and privacy-preserving data mining. In this section, we provide a more comprehensive definition of federated learning that considers data partitions, security, and applications. We also describe a workflow and system architecture for the federated-learning system.

### 2.1 Definition of Federated Learning

Define  $N$  data owners  $\{\mathcal{F}_1, \dots, \mathcal{F}_N\}$ , all of whom wish to train a machine-learning model by consolidating their respective data  $\{\mathcal{D}_1, \dots, \mathcal{D}_N\}$ . A conventional method is to put all data together and use  $\mathcal{D} = \mathcal{D}_1 \cup \dots \cup \mathcal{D}_N$  to train a model  $\mathcal{M}_{SUM}$ . A federated-learning system is a learning process in which the data owners collaboratively train a model  $\mathcal{M}_{FED}$ , in which process any data owner  $\mathcal{F}_i$  does not expose its data  $\mathcal{D}_i$  to others<sup>1</sup>. In addition, the accuracy of  $\mathcal{M}_{FED}$ , denoted as  $\mathcal{V}_{FED}$ , should be very close to the performance of  $\mathcal{M}_{SUM}$ ,  $\mathcal{V}_{SUM}$ . Formally, let  $\delta$  be a non-negative real number; if

$$|\mathcal{V}_{FED} - \mathcal{V}_{SUM}| < \delta, \quad (1)$$

we say that the federated learning algorithm has  $\delta$ -accuracy loss.

### 2.2 Privacy of Federated Learning

Privacy is one of the essential properties of federated learning. This requires security models and analysis to provide meaningful privacy guarantees. In this section, we briefly review and compare different privacy techniques for federated learning. We also identify approaches and potential challenges for preventing indirect leakage.

---

<sup>1</sup>The definition of data security may differ in different scenarios but is required to provide meaningful privacy guarantees. We demonstrate examples of security definitions in Section 2.3.

*Secure Multiparty Computation (SMC).* SMC security models involve multiple parties and provide security proof in a well-defined simulation framework to guarantee complete zero knowledge, that is, each party knows nothing except its input and output. Zero knowledge is highly desirable, but this desired property usually requires complicated computation protocols and may not be achieved efficiently. In certain scenarios, partial knowledge disclosure may be considered acceptable if security guarantees are provided. It is possible to build a security model with SMC under lower security requirements in exchange for efficiency [16]. Recently, a study [46] used the SMC framework for training machine-learning models with two servers and semi-honest assumptions. MPC protocols are used in [33] for model training and verification without users revealing sensitive data. One of the state-of-the-art SMC frameworks is Sharemind [8]. The authors of [44] proposed a 3PC model [5, 21, 45] with an honest majority and consider security in both semi-honest and malicious assumptions. These works require participants' data to be secretly shared among non-colluding servers.

*Differential Privacy.* Another line of work uses the techniques differential privacy [18] or k-anonymity [63] for data privacy protection [1, 12, 42, 61]. The methods of differential privacy, k-anonymity, and diversification [3] involve adding noise to the data, or using generalization methods to obscure certain sensitive attributes until the third party cannot distinguish the individual, thereby making the data impossible to be restore to protect user privacy. However, the root of these methods still requires that the data are transmitted elsewhere, which usually involves a trade-off between accuracy and privacy. In [23], the authors introduced a differential privacy approach to federated learning in order to add protection to client-side data by hiding the client's contributions during training.

*Homomorphic Encryption.* Homomorphic encryption [53] is also adopted to protect user data privacy through parameter exchange under the encryption mechanism during machine learning [24, 26, 48]. Unlike differential privacy protection, the data and the model itself are not transmitted, nor can they be guessed by the other party's data. Therefore, there is little possibility of leakage at the raw data level. Recent works adopted homomorphic encryption for centralizing and training data on the cloud [75, 76]. In practice, additively homomorphic encryption [2] is widely used and polynomial approximations need to be made to evaluate non-linear functions in machine-learning algorithms, resulting in trade-offs between accuracy and privacy [4, 35].

**2.2.1 Indirect Information Leakage.** Pioneer works of federated learning expose intermediate results such as parameter updates from an optimization algorithm like stochastic gradient descent (SGD) [41, 58]. However, no security guarantee is provided and the leakage of these gradients may actually leak important data information [51] when exposed together with data structure, such as in the case of image pixels. Researchers have considered the situation when one of the members of a federated learning system maliciously attacks others by allowing a back door to be inserted to learn others' data. In [6], the authors demonstrate that it is possible to insert hidden back doors into a joint global model and propose a new "constrain-and-scale" model-poisoning methodology to reduce data poisoning. In [43], researchers identified potential loopholes in collaborative machine-learning systems, where the training data used by different parties in collaborative learning is vulnerable to inference attacks. They showed that an adversarial participant can infer membership as well as properties associated with a subset of the training data. They also discussed possible defenses against these attacks. In [62], authors expose a potential security issue associated with gradient exchanges between different parties and propose a secured variant of the gradient descent method. They show that it tolerates up to a constant fraction of Byzantine workers.

Researchers have also started to consider a blockchain as a platform for facilitating federated learning. In [34], researchers have considered a blockchained federated-learning (BlockFL)

architecture, where mobile devices' local-learning model updates are exchanged and verified by leveraging a blockchain. They have considered an optimal block generation, network scalability, and robustness issues.

### 2.3 A Categorization of Federated Learning

In this section, we discuss how to categorize federated learning based on the distribution characteristics of the data.

Let matrix  $\mathcal{D}_i$  denote the data held by each data owner  $i$ . Each row of the matrix represents a sample, and each column represents a feature. At the same time, some datasets may also contain label data. We denote the feature space as  $\mathcal{X}$ , the label space as  $\mathcal{Y}$ , and we use  $\mathcal{I}$  to denote the sample ID space. For example, in the financial field, labels may be users' credit; in the marketing field, labels may be the user's purchase desire; in the education field,  $\mathcal{Y}$  may be the degree of the students. The feature  $\mathcal{X}$ , label  $\mathcal{Y}$ , and sample IDs  $\mathcal{I}$  constitute the complete training dataset  $(\mathcal{I}, \mathcal{X}, \mathcal{Y})$ . The feature and sample spaces of the data parties may not be identical, and we classify federated learning into horizontally federated learning, vertically federated learning, and federated transfer learning based on how data is distributed among various parties in the feature and sample ID space. Figure 2 shows the various federated learning frameworks for a two-party scenario.

**2.3.1 Horizontal Federated Learning.** Horizontal federated learning, or sample-based federated learning, is introduced in the scenarios in which datasets share the same feature space but different space in samples (Figure 2(a)). For example, two regional banks may have very different user groups from their respective regions, and the intersection set of their users is very small. However, their business is very similar, so the feature spaces are the same. The authors of [58] proposed a collaboratively deep-learning scheme in which participants train independently and share only subsets of updates of parameters. In 2017, Google proposed a horizontal federated-learning solution for Android phone model updates [41]. In that framework, a single user using an Android phone updates the model parameters locally and uploads the parameters to the Android cloud, thus jointly training the centralized model together with other data owners. A secure aggregation scheme to protect the privacy of aggregated user updates under their federated-learning framework is also introduced in [9]. The authors of [51] use additively homomorphic encryption for model parameter aggregation to provide security against the central server.

In [60], a multitask-style federated learning system is proposed to allow multiple sites to complete separate tasks while sharing knowledge and preserving security. Their proposed multitask learning model can also address high communication costs, stragglers, and fault tolerance issues. In [41], the authors proposed building a secure client–server structure in which the federated-learning system partitions data by users and allow models built at client devices to collaborate at the server site to build a global federated model. The process of model building ensures that there is no data leakage. Likewise, in [36], the authors proposed methods to improve the communication cost to facilitate the training of centralized models based on data distributed over mobile clients. Recently, a compression approach called deep gradient compression [39] was proposed to greatly reduce the communication bandwidth in large-scale distributed training.

We summarize horizontal federated learning as

$$\mathcal{X}_i = \mathcal{X}_j, \quad \mathcal{Y}_i = \mathcal{Y}_j, \quad I_i \neq I_j, \quad \forall \mathcal{D}_i, \mathcal{D}_j, i \neq j. \quad (2)$$

**Security Definition.** A horizontal federated learning system typically assumes honest participants and security against an honest-but-curious server [9, 51]. That is, only the server can compromise the privacy of data participants. Security proof has been provided in these works. Recently, another security model considering malicious users [29] was also proposed, posing additional privacy

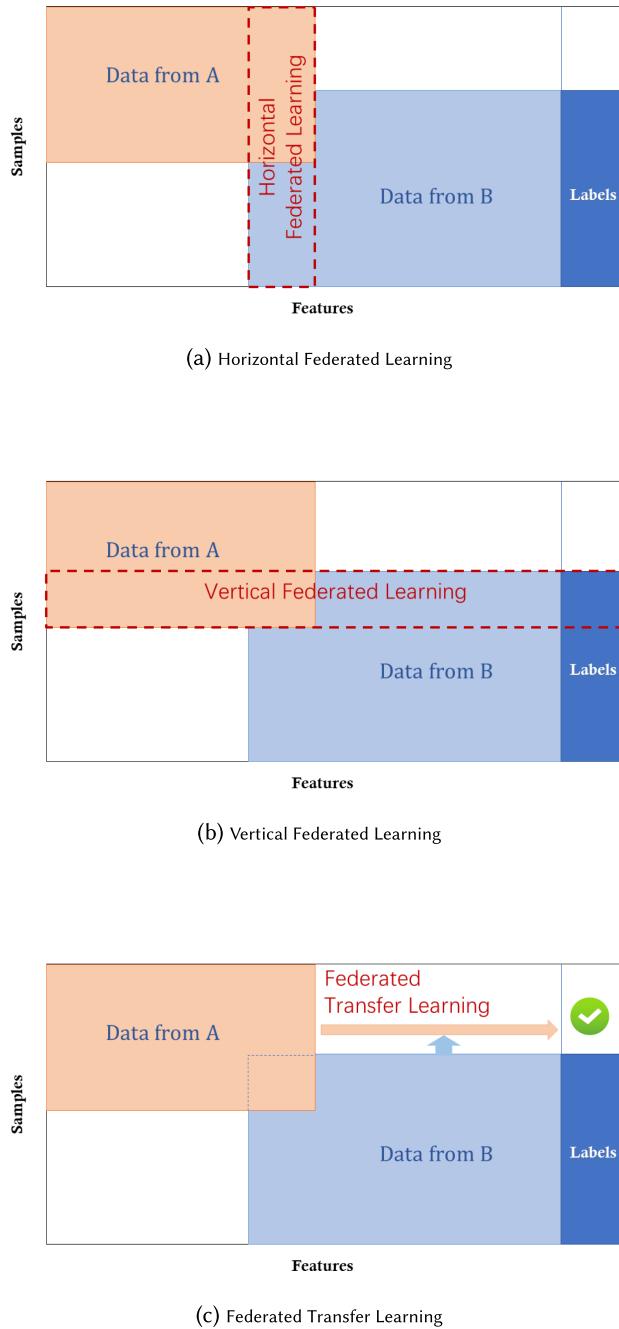


Fig. 2. Categorization of federated learning.

challenges. At the end of the training, the universal model and all of the model parameters are exposed to all participants.

**2.3.2 Vertical Federated Learning.** Privacy-preserving machine-learning algorithms have been proposed for vertically partitioned data, including cooperative statistical analysis [15], association rule mining [65], secure linear regression [22, 32, 55], classification [16], and gradient descent [68]. Recently, the authors of [27, 49] proposed a vertical federated-learning scheme to train a privacy-preserving logistic regression model. The authors studied the effect of entity resolution on learning performance and applied Taylor approximation to the loss and gradient functions so that homomorphic encryption can be adopted for privacy-preserving computations.

Vertical federated learning or feature-based federated learning (Figure 2(b)) is applicable to the cases in which two datasets share the same sample ID space but differ in feature space. For example, consider two different companies in the same city: one is a bank and the other is an e-commerce company. Their user sets are likely to contain most of the residents of the area; thus, the intersection of their user space is large. However, since the bank records the user's revenue and expenditure behavior and credit rating and the e-commerce retains the user's browsing and purchasing history, their feature spaces are very different. Suppose that we want both parties to have a prediction model for product purchases based on user and product information.

Vertically federated learning is the process of aggregating these different features and computing the training loss and gradients in a privacy-preserving manner to build a model with data from both parties collaboratively. Under such a federal mechanism, the identity and the status of each participating party is the same, and the federal system helps everyone establish a “common wealth” strategy, which is why this system is called federated learning. Therefore, in such a system, we have that

$$\mathcal{X}_i \neq \mathcal{X}_j, \quad \mathcal{Y}_i \neq \mathcal{Y}_j, \quad I_i = I_j \quad \forall \mathcal{D}_i, \mathcal{D}_j, i \neq j. \quad (3)$$

**Security Definition.** A vertical federated-learning system typically assumes honest but curious participants. In a two-party case, for example, the two parties are non-colluding and at most one is compromised by an adversary. The security definition is that the adversary can learn data only from the client that it corrupted and not data from the other client beyond what is revealed by the input and output. To facilitate secure computations between the two parties, sometimes a semi-honest third party (STP) is introduced, in which case it is assumed that the STP does not collude with either party. SMC provides formal privacy proof for these protocols [25]. At the end of learning, each party holds only those model parameters associated to its own features. Therefore, at inference time, the two parties also need to collaborate to generate output.

**2.3.3 Federated Transfer Learning (FTL).** Federated transfer learning applies to the scenarios in which two datasets differ not only in samples but also in feature space. Consider two institutions: one is a bank located in China and the other is an e-commerce company located in the United States. Owing to geographical restrictions, the user groups of the two institutions have a small intersection. On the other hand, owing to the different businesses, only a small portion of the feature space from both parties overlaps. In this case, transfer-learning [50] techniques can be applied to provide solutions for the entire sample and feature space under a federation (Figure 2(c)). Specifically, a common representation between the two feature spaces is learned using limited common sample sets and later applied to obtain predictions for samples with only one-side features. FTL is an important extension to the existing federated learning systems because it deals with problems exceeding the scope of existing federated learning algorithms:

$$\mathcal{X}_i \neq \mathcal{X}_j, \quad \mathcal{Y}_i \neq \mathcal{Y}_j, \quad I_i \neq I_j \quad \forall \mathcal{D}_i, \mathcal{D}_j, i \neq j. \quad (4)$$

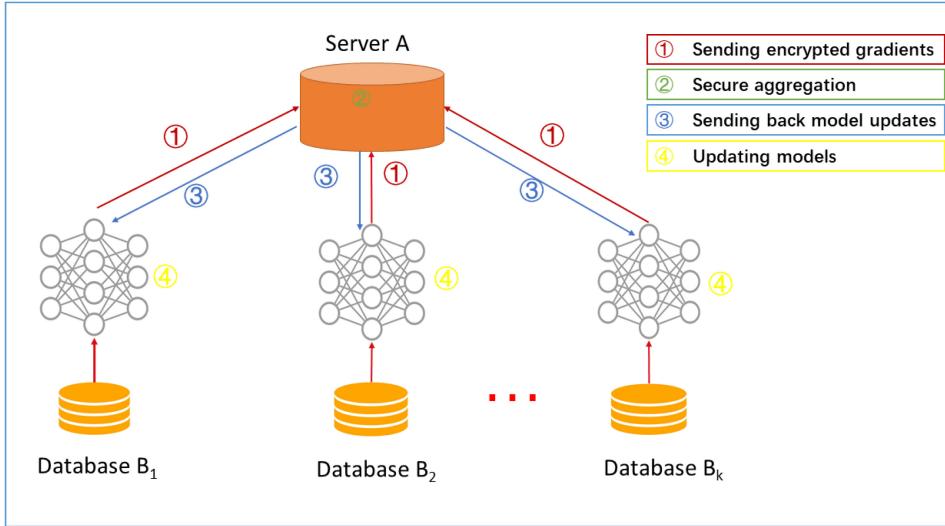


Fig. 3. Architecture for a horizontal federated-learning system.

*Security Definition.* A federated transfer learning system typically involves two parties. As will be shown in the next section, its protocols are similar to the ones in vertical federated learning, in which case the security definition for vertical federated learning can be extended here.

#### 2.4 Architecture for a Federated Learning System

In this section, we illustrate examples of general architectures for a federated-learning system. Note that the architectures of horizontal and vertical federated-learning systems are quite different by design, and we will introduce them separately.

**2.4.1 Horizontal Federated Learning.** A typical architecture for a horizontal federated-learning system is shown in Figure 3. In this system,  $k$  participants with the same data structure collaboratively learn a machine-learning model with the help of a parameter or cloud server. A typical assumption is that the participants are honest whereas the server is honest but curious; therefore, no leakage of information from any participants to the server is allowed [51]. The training process of such a system usually contains the following four steps.

- **Step 1:** Participants locally compute training gradients; mask a selection of gradients with encryption [51], differential privacy [58], or secret sharing [9] techniques; and send masked results to the server.
- **Step 2:** The server performs secure aggregation without learning information about any participant.
- **Step 3:** The server sends back the aggregated results to participants.
- **Step 4:** Participants update their respective model with the decrypted gradients.

Iterations through the above steps continue until the loss function converges, thus completing the entire training process. This architecture is independent of specific machine-learning algorithms (logistic regression, DNN, etc.) and all participants will share the final model parameters.

*Security Analysis.* The above architecture is proved to protect data leakage against the semi-honest server if gradient aggregation is done with SMC [9] or homomorphic encryption [51].

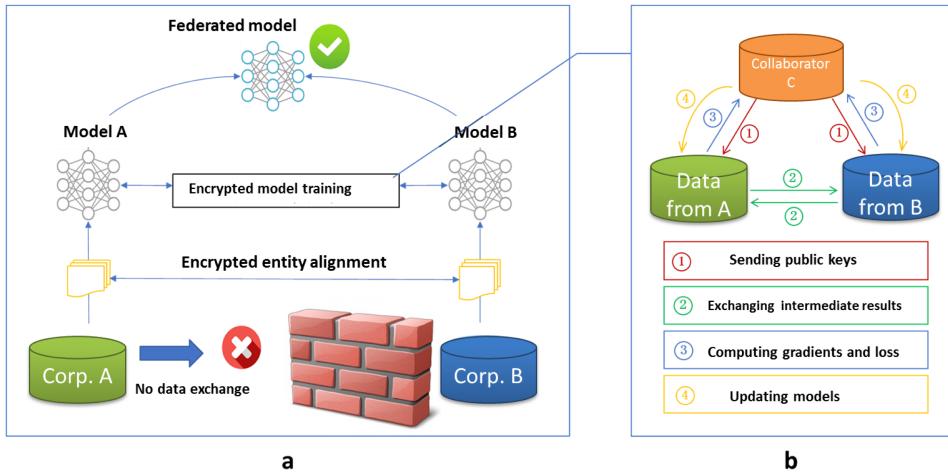


Fig. 4. Architecture for a vertical federated-learning system.

However, it may be subject to attack in another security model by a malicious participant training a generative adversarial network (GAN) in the collaborative-learning process [29].

**2.4.2 Vertical Federated Learning.** Suppose that companies A and B would like to jointly train a machine-learning model and their business systems each have their own data. In addition, Company B also has label data that the model needs to predict. For data privacy and security reasons, A and B cannot directly exchange data. In order to ensure the confidentiality of the data during the training process, a third-party collaborator C is involved. Here, we assume that the collaborator C is honest and does not collude with A or B, but party A and B are honest but curious to each other. A trusted third-party C is a reasonable assumption since party C can be played by authorities such as governments or replaced by a secure computing node such as Intel Software Guard Extensions (SGXs) [7]. The federated-learning system consists of two parts, as shown in Figure 4.

*Part 1.* Encrypted entity alignment. Since the user groups of the two companies are not the same, the system uses the encryption-based user ID alignment techniques such as those described in [38, 56] to confirm the common users of both parties without A and B exposing their respective data. During the entity alignment, the system does not expose users that do not overlap with each other.

*Part 2.* Encrypted model training. After determining the common entities, we can use these common entities' data to train the machine-learning model. The training process can be divided into the following four steps (as shown in Figure 4).

- **Step 1:** Collaborator C creates encryption pairs and sends a public key to A and B.
- **Step 2:** A and B encrypt and exchange the intermediate results for gradient and loss calculations.
- **Step 3:** A and B compute encrypted gradients and add an *additional mask*, respectively. B also computes encrypted loss. A and B send encrypted values to C.
- **Step 4:** C decrypts and send the decrypted gradients and loss back to A and B. A and B unmask the gradients and update the model parameters accordingly.

Here, we illustrate the training process using linear regression and homomorphic encryption as an example. To train a linear regression model with gradient descent methods, we need secure computations of its loss and gradients. Assuming learning rate  $\eta$ , regularization parameter  $\lambda$ , dataset

Table 1. Training Steps for Vertical Federated Learning: Linear Regression

	Party A	Party B	Party C
Step 1	Initializes $\Theta_A$ .	Initializes $\Theta_B$ .	Creates an encryption key pair, sends public key to A and B.
Step 2	Compute $[[u_i^A]]$ , $[[\mathcal{L}_{AB}]]$ and sends to B.	Compute $[[u_i^B]]$ , $[[d_i^B]]$ , $[[\mathcal{L}]]$ , sends $[[d_i^B]]$ to A, sends $[[\mathcal{L}]]$ to C.	
Step 3	Initializes $R_A$ , compute $[[\frac{\partial \mathcal{L}}{\partial \Theta_A}]] + [[R_A]]$ and sends to C.	Initializes $R_B$ , compute $[[\frac{\partial \mathcal{L}}{\partial \Theta_B}]] + [[R_B]]$ and sends to C.	Decrypts $\mathcal{L}$ , sends $\frac{\partial \mathcal{L}}{\partial \Theta_A} + R_A$ to A, $\frac{\partial \mathcal{L}}{\partial \Theta_B} + R_B$ to B.
Step 4	Updates $\Theta_A$ .	Updates $\Theta_B$ .	
What is obtained	$\Theta_A$	$\Theta_B$	

Table 2. Evaluation Steps for Vertical Federated Learning : Linear Regression

	Party A	Party B	Inquisitor
Step 0			Sends user ID $i$ to A and B.
Step 1	Compute $u_i^A$ and sends to C.	Compute $u_i^B$ and sends to C.	Gets result $u_i^A + u_i^B$ .

$\{x_i^A\}_{i \in \mathcal{D}_A}$ ,  $\{x_i^B, y_i\}_{i \in \mathcal{D}_B}$ , and model parameters  $\Theta_A$ ,  $\Theta_B$  corresponding to the feature space of  $x_i^A$ ,  $x_i^B$  respectively, the training objective is

$$\min_{\Theta_A, \Theta_B} \sum_i \|\Theta_A x_i^A + \Theta_B x_i^B - y_i\|^2 + \frac{\lambda}{2} (\|\Theta_A\|^2 + \|\Theta_B\|^2). \quad (5)$$

Let  $u_i^A = \Theta_A x_i^A$ ,  $u_i^B = \Theta_B x_i^B$ ; then, the encrypted loss is

$$[[\mathcal{L}]] = \left[ \left[ \sum_i ((u_i^A + u_i^B - y_i)^2) + \frac{\lambda}{2} (\Theta_A^2 + \Theta_B^2) \right] \right], \quad (6)$$

where additive homomorphic encryption is denoted as  $[[\cdot]]$ . Let  $[[\mathcal{L}_A]] = [[\sum_i ((u_i^A)^2) + \frac{\lambda}{2} \Theta_A^2]]$ ,  $[[\mathcal{L}_B]] = [[\sum_i ((u_i^B - y_i)^2) + \frac{\lambda}{2} \Theta_B^2]]$ , and  $[[\mathcal{L}_{AB}]] = 2 \sum_i ([[u_i^A]] (u_i^B - y_i))$ ; then,

$$[[\mathcal{L}]] = [[\mathcal{L}_A]] + [[\mathcal{L}_B]] + [[\mathcal{L}_{AB}]]. \quad (7)$$

Similarly, let  $[[d_i]] = [[u_i^A]] + [[u_i^B - y_i]]$ ; then, the gradients are

$$\left[ \left[ \frac{\partial \mathcal{L}}{\partial \Theta_A} \right] \right] = \sum_i [[d_i]] x_i^A + [[\lambda \Theta_A]] \quad (8)$$

$$\left[ \left[ \frac{\partial \mathcal{L}}{\partial \Theta_B} \right] \right] = \sum_i [[d_i]] x_i^B + [[\lambda \Theta_B]]. \quad (9)$$

See Tables 1 and 2 for the detailed steps. During entity alignment and model training, the data of A and B are kept locally and the data interaction in training does not lead to data privacy leakage. Note that potential information leakage to C may or may not be considered to be privacy violation. To further prevent C from learning information from A or B in this case, A and B can further hide

their gradients from C by adding encrypted random masks. Therefore, the two parties achieve training of a common model cooperatively with the help of federated learning. Because, during the training, the loss and gradients that each party receives are exactly the same as the loss and gradients they would receive, if jointly building a model with data gathered at one place without privacy constraints, that is, this model is lossless. The efficiency of the model depends on the communication cost and computation cost of encrypted data. In each iteration, the information sent between A and B scales with the number of overlapping samples. Therefore, the efficiency of this algorithm can be further improved by adopting distributed parallel computing techniques.

*Security Analysis.* The training protocol shown in Table 1 does not reveal any information to C, because all that C learns are the masked gradients, and the randomness and secrecy of the masked matrix are guaranteed [16]. In the above protocol, party A learns its gradient at each step, but this is not enough for A to learn any information from B according to Equation (8), because the security of scalar product protocol is well established based on the inability of solving n equations in more than n unknowns [16, 65]. Here, we assume that the number of samples  $N_A$  is much greater than  $n_A$ , where  $n_A$  is the number of features. Similarly, party B cannot learn any information from A. Therefore, the security of the protocol is proved. Note that we have assumed that both parties are semi-honest. If a party is malicious and cheats the system by faking its input—for example, party A submits only one non-zero input with only one non-zero feature—it can tell the value of  $u_i^B$  for that feature of that sample. It still cannot tell  $x_i^B$  or  $\Theta_B$ , though, and the deviation will distort results for the next iteration, alarming the other party, who will terminate the learning process. At the end of the training process, each party (A or B) remains oblivious to the data structure of the other party and obtains the model parameters associated only with its own features. At inference time, the two parties need to collaboratively compute the prediction results with the steps shown in Table 2, which still do not lead to information leakage.

**2.4.3 Federated Transfer Learning.** Suppose that in the above vertical federated-learning example, parties A and B have only a very small set of overlapping samples and we are interested in learning the labels for the entire dataset in party A. The architecture described in the above section so far works only for the overlapping dataset. To extend its coverage to the entire sample space, we introduce transfer learning. This does not change the overall architecture shown in Figure 4, but does change the details of the intermediate results that are exchanged between parties A and B. Specifically, transfer learning typically involves learning a common representation between the features of parties A and B and minimizing the errors in predicting the labels for the target-domain party by leveraging the labels in the source-domain party (B in this case). Therefore, the gradient computations for parties A and B are different from that in the vertical federated-learning scenario. At inference time, it still requires both parties to compute the prediction results.

**2.4.4 Incentives Mechanism.** In order to fully commercialize federated learning among different organizations, a fair platform and incentive mechanisms needs to be developed [20]. After the model is built, the performance of the model will be manifested in the actual applications. This performance can be recorded in a permanent data-recording mechanism (such as Blockchain). Organizations that provide more data will be better off, and the model's effectiveness depends on the data provider's contribution to the system. The effectiveness of these models is distributed to parties based on federated mechanisms and continue to motivate more organizations to join the data federation.

The implementation of the above architecture not only considers the privacy protection and effectiveness of collaborative modeling among multiple organizations but also considers how to

reward organizations that contribute more data, and how to implement incentives with a consensus mechanism. Therefore, federated learning is a “closed-loop” learning mechanism.

### 3 RELATED WORKS

Federated learning enables multiple parties to collaboratively construct a machine-learning model while keeping their private training data private. As a novel technology, federated learning has several threads of originality, some of which are rooted on existing fields. Below, we explain the relationship between federated learning and other related concepts from multiple perspectives.

#### 3.1 Privacy-Preserving Machine Learning

Federated learning can be considered as privacy-preserving, decentralized collaborative machine learning. Therefore, it is tightly related to multiparty, privacy-preserving machine learning. Many research efforts have been devoted to this area in the past. For example, the authors of [17, 67] proposed algorithms for secure multiparty decision trees for vertically partitioned data. Vaidya and Clifton proposed secure association mining rules [65], secure k-means [66], and a naïve Bayes classifier [64] for vertically partitioned data. The authors of [31] proposed an algorithm for association rules on horizontally partitioned data. Secure support vector machines algorithms have been developed for vertically partitioned data [73] and horizontally partitioned data [74]. The authors of [16] proposed secure protocols for multiparty linear regression and classification. The authors of [68] proposed secure multiparty gradient descent methods. These works all used SMC [25, 72] for privacy guarantees.

Nikolaenko et al. [48] implemented a privacy-preserving protocol for linear regression on horizontally partitioned data using homomorphic encryption and Yao’s garbled circuits. The authors of [22, 24] proposed a linear regression approach for vertically partitioned data. These systems solved the linear regression problem directly. The authors of [47] approached the problem with SGD and also proposed privacy-preserving protocols for logistic regression and neural networks. Recently, a follow-up work with a three-server model was proposed [44]. Aono et al. [4] proposed a secure logistic regression protocol using homomorphic encryption. Shokri and Shmatikov [58] proposed training of neural networks for horizontally partitioned data with exchanges of updated parameters. The authors of [51] used additively homomorphic encryption to preserve the privacy of gradients and enhance the security of the system. With the recent advances in deep learning, privacy-preserving neural networks inference is also receiving a lot of research interest [10, 11, 14, 28, 40, 52, 54].

#### 3.2 Federated Learning versus Distributed Machine Learning

Horizontal federated learning at first sight is somewhat similar to distributed machine learning. Distributed machine learning covers many aspects, including distributed storage of training data, distributed operation of computing tasks, and distributed distribution of model results. A parameter server [30] is a typical element in distributed machine learning. As a tool to accelerate the training process, the parameter server stores data on distributed working nodes and allocates data and computing resources through a central scheduling node to train the model more efficiently. For horizontally federated learning, the working node represents the data owner. It has full autonomy for the local data; it can decide when and how to join the federated learning. In the parameter server, the central node always takes control; thus, federated learning is faced with a more complex learning environment. In addition, federated learning emphasizes the data-privacy protection of the data owner during the model training process. Effective measures to protect data privacy can better cope with the increasingly stringent data privacy and data security regulatory environment in the future.

As in distributed machine-learning settings, federated learning will also need to address non-IID data. The authors of [77] showed that, with non-IID local data, performance can be greatly reduced for federated learning. The authors in response supplied a new method to address the issue similar to transfer learning.

### 3.3 Federated Learning versus Edge Computing

Federated learning can be seen as an operating system for edge computing, as it provides the learning protocol for coordination and security. The authors of [69] considered a generic class of machine-learning models that are trained using gradient descent-based approaches. They analyze the convergence bound of distributed gradient descent from a theoretical point of view, based on which they propose a control algorithm that determines the best trade-off between local update and global parameter aggregation to minimize the loss function under a given resource budget.

### 3.4 Federated Learning versus Federated Database Systems

Federated database systems [57] are systems that integrate multiple database units and manage the integrated system as a whole. The federated database concept is proposed to achieve interoperability with multiple independent databases. A federated database system often uses distributed storage for database units and, in practice, the data in each database unit is heterogeneous. Therefore, it has many similarities with federated learning in terms of the type and storage of data. However, the federated database system does not involve any privacy protection mechanism in the process of interacting with each other, and all database units are completely visible to the management system. In addition, the focus of the federated database system is on the basic operations of data—including inserting, deleting, searching, and merging—while the purpose of federated learning is to establish a joint model for each data owner under the premise of protecting data privacy, so that the various values and laws that the data contain serve us better.

## 4 APPLICATIONS

As an innovative modeling mechanism that could train a united model on data from multiple parties without compromising privacy and security of those data, federated learning has a promising application in sales, financial, and many other industries in which data cannot be directly aggregated for training machine-learning models owing to factors such as intellectual property rights, privacy protection, and data security.

Take smart retail as an example. Its purpose is to use machine-learning techniques to provide customers with personalized services, mainly including product recommendation and sales services. The data features involved in the smart retail business mainly include user purchasing power, user personal preference, and product characteristics. In practical applications, these three data features are likely to be scattered among three different departments or enterprises. For example, a user's purchasing power can be inferred from the user's bank savings and personal preference can be analyzed from the user's social networks, while the characteristics of products are recorded by an e-shop. In this scenario, we are facing two problems. First, for the protection of data privacy and data security, data barriers between banks, social networking sites, and e-shopping sites are difficult to break. As a result, data cannot be directly aggregated to train a model. Second, the data stored by the three parties are usually heterogeneous, and traditional machine-learning models cannot directly work on heterogeneous data. For now, these problems have not been effectively solved with traditional machine-learning methods, which hinder the popularization and application of AI in more fields.

Federated learning and transfer learning are the key to solving these problems. First, by exploiting the characteristics of federated learning, we can build a machine-learning model for the three

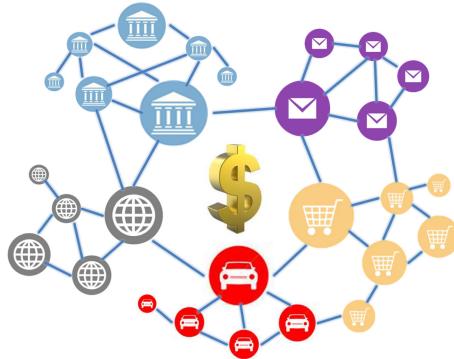


Fig. 5. Data alliance allocates the benefits on a blockchain.

parties without exporting the enterprise data, which not only fully protects data privacy and data security but also provides customers with personalized and targeted services and thereby achieves mutual benefits. Meanwhile, we can leverage transfer learning to address the data heterogeneity problem and break through the limitations of traditional AI techniques. Therefore, federated learning provides good technical support for us to build a cross-enterprise, cross-data, and cross-domain ecosphere for big data and AI.

One can use the federated-learning framework for multiparty database querying without exposing the data. For example, suppose that in a finance application we are interested in detecting multiparty borrowing, which has been a major risk factor in the banking industry. This happens when certain users maliciously borrow from one bank to pay for the loan at another bank. Multiparty borrowing is a threat to financial stability, as a large number of such illegal actions may cause the entire financial system to collapse. To find such users without exposing the user lists to each other between banks  $A$  and  $B$ , we can exploit a federated-learning framework. In particular, we can use the encryption mechanism of federated learning and encrypt the user list at each party and then take the intersection of the encrypted list in the federation. The decryption of the final result gives the list of multiparty borrowers without exposing the other “good” users to the other party. As we will see below, this operation corresponds to the vertical federated learning framework.

Smart healthcare is another domain that we expect will greatly benefit from the rising of federated-learning techniques. Medical data such as disease symptoms, gene sequences, and medical reports are very sensitive and private, yet medical datasets are difficult to collect and exist in isolated medical centers and hospitals. The insufficiency of data sources and the lack of labels have led to an unsatisfactory performance of machine-learning models, which has become the bottleneck of current smart healthcare. We envisage that if all medical institutions are united and share their data to form a large medical dataset, then the performance of machine-learning models trained on that large medical dataset would be significantly improved. Federated learning combining with transfer learning is the main way to achieve this vision. Transfer learning could be applied to fill the missing labels, thereby expanding the scale of the available data and further improving the performance of a trained model. Therefore, federated transfer learning would play a pivotal role in the development of smart healthcare and may be able to take human healthcare to a whole new level.

## 5 FEDERATED LEARNING AND DATA ALLIANCE OF ENTERPRISES

Federated learning is not only a technology standard but also a business model. When people realize the effects of big data, the first thought that occurs to them is to aggregate the data, compute

the models through a remote processor, and then download the results for further use. Cloud computing comes into being under such demands. That being said, with the increasing importance of data privacy and data security and a closer relationship between a company's profits and its data, the cloud computing model has been challenged. However, the business model of federated learning has provided a new paradigm for applications of big data. When the isolated data occupied by each institution fails to produce an ideal model, the mechanism of federated learning makes it possible for institutions and enterprises to share a united model without data exchange. Furthermore, federated learning could make equitable rules for profits allocation with the help of a consensus mechanism from blockchain techniques. The data possessors, regardless of the scale of data that they have, will be motivated to join in the data alliance and make their own profits. We believe that the establishment of the business model for data alliance and the technical mechanism for federated learning should be carried out together. We would also make standards for federated learning in various fields to put it into use as soon as possible.

## 6 CONCLUSIONS AND PROSPECTS

Recently, the isolation of data and the emphasis on data privacy became the next challenges for AI, but federated learning has brought us new hope. It could establish a united model for multiple enterprises while local data is protected so that enterprises could work together on data security. This article generally introduces the basic concept, architecture, and techniques of federated learning, and discusses its potential in various applications. It is expected that, in the near future, federated learning would break the barriers between industries and establish a community where data and knowledge could be shared with safety and the benefits would be fairly distributed according to the contribution of each participant. The bonus of AI would finally be brought to every corner of our lives.

## REFERENCES

- [1] Martin Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. 2016. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS'16)*. ACM, New York, NY, 308–318. DOI : <https://doi.org/10.1145/2976749.2978318>
- [2] Abbas Acar, Hidayet Aksu, A. Selcuk Uluagac, and Mauro Conti. 2018. A survey on homomorphic encryption schemes: Theory and implementation. *ACM Comput. Surv.* 51, 4, Article 79 (July 2018), 35 pages. DOI : <https://doi.org/10.1145/3214303>
- [3] Rakesh Agrawal and Ramakrishnan Srikant. 2000. Privacy-preserving data mining. In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data (SIGMOD'00)*. ACM, New York, NY, 439–450. DOI : <https://doi.org/10.1145/342009.335438>
- [4] Yoshinori Aono, Takuya Hayashi, Le Trieu Phong, and Lihua Wang. 2016. Scalable and secure logistic regression via homomorphic encryption. In *Proceedings of the 6th ACM Conference on Data and Application Security and Privacy (CODASPY'16)*. ACM, New York, NY, 142–144. DOI : <https://doi.org/10.1145/2857705.2857731>
- [5] Toshinori Araki, Jun Furukawa, Yehuda Lindell, Ariel Nof, and Kazuma Ohara. 2016. High-throughput semi-honest secure three-party computation with an honest majority. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS'16)*. ACM, New York, NY, 805–817. DOI : <https://doi.org/10.1145/2976749.2978331>
- [6] Eugene Bagdasaryan, Andreas Veit, Yiqing Hua, Deborah Estrin, and Vitaly Shmatikov. 2018. How To Backdoor Federated Learning. arxiv:cs.CR/1807.00459
- [7] Raad Bahmani, Manuel Barbosa, Ferdinand Brassier, Bernardo Portela, Ahmad-Reza Sadeghi, Guillaume Scerri, and Bogdan Warinschi. 2017. Secure multiparty computation from SGX. *Financial Cryptography*. 477–497.
- [8] Dan Bogdanov, Sven Laur, and Jan Willemson. 2008. Sharemind: A framework for fast privacy-preserving computations. In *Proceedings of the 13th European Symposium on Research in Computer Security: Computer Security (ESORICS'08)*. Springer, Berlin, 192–206. DOI : [https://doi.org/10.1007/978-3-540-88313-5\\_13](https://doi.org/10.1007/978-3-540-88313-5_13)
- [9] Keith Bonawitz, Vladimir Ivanov, Ben Kreuter, Antonio Marcedone, H. Brendan McMahan, Sarvar Patel, Daniel Ramage, Aaron Segal, and Karn Seth. 2017. Practical secure aggregation for privacy-preserving machine learning. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security (CCS'17)*. ACM, New York, NY, 1175–1191. DOI : <https://doi.org/10.1145/3133956.3133982>

- [10] Florian Bourse, Michele Minelli, Matthias Minihold, and Pascal Paillier. 2017. Fast homomorphic evaluation of deep discretized neural networks. *IACR Cryptology ePrint Archive* 2017 (2017), 1114.
- [11] Hervé Chabanne, Amaury de Wargny, Jonathan Milgram, Constance Morel, and Emmanuel Prouff. 2017. Privacy-preserving classification on deep neural network. *IACR Cryptology ePrint Archive* 2017 (2017), 35.
- [12] Kamalika Chaudhuri and Claire Monteleoni. 2009. Privacy-preserving logistic regression. In *Advances in Neural Information Processing Systems 21*, D. Koller, D. Schuurmans, Y. Bengio, and L. Bottou (Eds.). Curran Associates, Inc., 289–296. <http://papers.nips.cc/paper/3486-privacy-preserving-logistic-regression.pdf>.
- [13] Fei Chen, Zhenhua Dong, Zhenguo Li, and Xiuqiang He. 2018. Federated meta-learning for recommendation. *CoRR* abs/1802.07876 (2018). arxiv:1802.07876 <http://arxiv.org/abs/1802.07876>.
- [14] Nathan Dowlin, Ran Gilad-Bachrach, Kim Laine, Kristin Lauter, Michael Naehrig, and John Wernsing. 2016. *CryptoNets: Applying Neural Networks to Encrypted Data with High Throughput and Accuracy*. Technical Report. Retrieved December 26, 2018 from <https://www.microsoft.com/en-us/research/publication/cryptonets-applying-neural-networks-to-encrypted-data-with-high-throughput-and-accuracy/>.
- [15] W. Du and M. Atallah. 2001. Privacy-preserving cooperative statistical analysis. In *Proceedings of the 17th Annual Computer Security Applications Conference (ACSAC'01)*. IEEE Computer Society, Washington, DC, 102–. <http://dl.acm.org/citation.cfm?id=872016.872181>.
- [16] Wenliang Du, Yunhsiang Sam Han, and Shigang Chen. 2004. Privacy-preserving multivariate statistical analysis: Linear regression and classification. In *SDM*, Vol. 4. 222–233.
- [17] Wenliang Du and Zhijun Zhan. 2002. Building decision tree classifier on private data. In *Proceedings of the IEEE International Conference on Privacy, Security and Data Mining - Volume 14 (CRPIT'02)*. Australian Computer Society, Inc., Darlinghurst, Australia, Australia, 1–8. <http://dl.acm.org/citation.cfm?id=850782.850784>.
- [18] Cynthia Dwork. 2008. Differential privacy: A survey of results. In *Proceedings of the 5th International Conference on Theory and Applications of Models of Computation (TAMC'08)*. Springer, Berlin, 1–19. <http://dl.acm.org/citation.cfm?id=1791834.1791836>.
- [19] EU. 2016. REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing directive 95/46/EC (general data protection regulation). Retrieved December 26, 2018 from <https://eur-lex.europa.eu/legal-content/EN/TXT>.
- [20] Boi Faltings, Goran Radanovic, and Ronald Brachman. 2017. *Game Theory for Data Science: Eliciting Truthful Information*. Morgan & Claypool Publishers.
- [21] Jun Furukawa, Yehuda Lindell, Ariel Nof, and Or Weinstein. 2016. High-Throughput Secure Three-Party Computation for Malicious Adversaries and an Honest Majority. *Cryptology ePrint Archive*, Report 2016/944. <https://eprint.iacr.org/2016/944>.
- [22] Adrià Gascón, Philipp Schoppmann, Borja Balle, Mariana Raykova, Jack Doerner, Samee Zahur, and David Evans. 2016. Secure linear regression on vertically partitioned datasets. *IACR Cryptology ePrint Archive* 2016 (2016), 892.
- [23] Robin C. Geyer, Tassilo Klein, and Moin Nabi. 2017. Differentially private federated learning: A client level perspective. *CoRR* abs/1712.07557 (2017). arxiv:1712.07557 <http://arxiv.org/abs/1712.07557>
- [24] Irene Giacomelli, Somesh Jha, Marc Joye, C. David Page, and Kyonghwan Yoon. 2017. Privacy-preserving ridge regression with only linearly-homomorphic encryption. *Cryptology ePrint Archive, Report* 2017/979. <https://eprint.iacr.org/2017/979>.
- [25] O. Goldreich, S. Micali, and A. Wigderson. 1987. How to play any mental game. In *Proceedings of the 19th Annual ACM Symposium on Theory of Computing (STOC'87)*. ACM, New York, NY, 218–229. DOI: <https://doi.org/10.1145/28395.28420>
- [26] Rob Hall, Stephen E. Fienberg, and Yuval Nardi. 2011. Secure multiple linear regression based on homomorphic encryption. *Journal of Official Statistics* 27, 4 (2011), 669–691.
- [27] Stephen Hardy, Wilko Henecka, Hamish Ivey-Law, Richard Nock, Giorgio Patrini, Guillaume Smith, and Brian Thorne. 2017. Private federated learning on vertically partitioned data via entity resolution and additively homomorphic encryption. *CoRR* abs/1711.10677 (2017).
- [28] Ehsan Hesamifard, Hassan Takabi, and Mehdi Ghasemi. 2017. CryptoDL: Deep neural networks over encrypted data. *CoRR* abs/1711.05189 (2017). arxiv:1711.05189 <http://arxiv.org/abs/1711.05189>.
- [29] Briland Hitaj, Giuseppe Ateniese, and Fernando Pérez-Cruz. 2017. Deep models under the GAN: Information leakage from collaborative deep learning. *CoRR* abs/1702.07464 (2017).
- [30] Qirong Ho, James Cipar, Henggang Cui, Jin Kyu Kim, Seunghak Lee, Phillip B. Gibbons, Garth A. Gibson, Gregory R. Ganger, and Eric P. Xing. 2013. More effective distributed ML via a stale synchronous parallel parameter server. In *Proceedings of the 26th International Conference on Neural Information Processing Systems - Volume 1 (NIPS'13)*. Curran Associates Inc., 1223–1231. <http://dl.acm.org/citation.cfm?id=2999611.2999748>.

- [31] Murat Kantarcioglu and Chris Clifton. 2004. Privacy-preserving distributed mining of association rules on horizontally partitioned data. *IEEE Trans. on Knowl. and Data Eng.* 16, 9 (Sept. 2004), 1026–1037. DOI : <https://doi.org/10.1109/TKDE.2004.45>
- [32] Alan F. Karr, X. Sheldon Lin, Ashish P. Sanil, and Jerome P. Reiter. 2004. Privacy-preserving analysis of vertically partitioned data using secure matrix products. *Journal of Official Statistics* 25, 125–138.
- [33] Niki Kilbertus, Adria Gascon, Matt Kusner, Michael Veale, Krishna Gummadi, and Adrian Weller. 2018. Blind justice: Fairness with encrypted sensitive attributes. In *Proceedings of the 35th International Conference on Machine Learning (Proceedings of Machine Learning Research)*, Jennifer Dy and Andreas Krause (Eds.), Vol. 80. PMLR, Stockholm, Sweden, 2630–2639. <http://proceedings.mlr.press/v80/kilbertus18a.html>.
- [34] Hyesung Kim, Jihong Park, Mehdi Bennis, and Seong-Lyun Kim. 2018. On-Device Federated Learning via Blockchain and its Latency Analysis. arxiv:cs.IT/1808.03949
- [35] Miran Kim, Yongsoo Song, Shuang Wang, Yuhou Xia, and Xiaoqian Jiang. 2018. Secure logistic regression based on homomorphic encryption: Design and evaluation. *JMIR Med Inform* 6, 2 (17 Apr 2018), e19. DOI : <https://doi.org/10.2196/medinform.8805>
- [36] Jakub Konecný, H. Brendan McMahan, Daniel Ramage, and Peter Richtárik. 2016. Federated optimization: Distributed machine learning for on-device intelligence. *CoRR* abs/1610.02527 (2016). arxiv:1610.02527 <http://arxiv.org/abs/1610.02527>
- [37] Jakub Konecný, H. Brendan McMahan, Felix X. Yu, Peter Richtárik, Ananda Theertha Suresh, and Dave Bacon. 2016. Federated learning: Strategies for improving communication efficiency. *CoRR* abs/1610.05492 (2016). arxiv:1610.05492 <http://arxiv.org/abs/1610.05492>
- [38] Gang Liang and Sudarshan S. Chawathe. 2004. Privacy-preserving inter-database operations. In *International Conference on Intelligence and Security Informatics*. Springer, 66–82.
- [39] Yujun Lin, Song Han, Huizi Mao, Yu Wang, and William J. Dally. 2017. Deep gradient compression: Reducing the communication bandwidth for distributed training. *CoRR* abs/1712.01887 (2017). arxiv:1712.01887 <http://arxiv.org/abs/1712.01887>.
- [40] Jian Liu, Mika Juuti, Yao Lu, and N. Asokan. 2017. Oblivious neural network predictions via MiniONN transformations. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security (CCS’17)*. ACM, New York, NY, 619–631. DOI : <https://doi.org/10.1145/3133956.3134056>
- [41] H. Brendan McMahan, Eider Moore, Daniel Ramage, and Blaise Agüera y Arcas. 2016. Federated learning of deep networks using model averaging. *CoRR* abs/1602.05629 (2016). arxiv:1602.05629 <http://arxiv.org/abs/1602.05629>.
- [42] H. Brendan McMahan, Daniel Ramage, Kunal Talwar, and Li Zhang. 2017. Learning differentially private language models without losing accuracy. *CoRR* abs/1710.06963 (2017).
- [43] Luca Melis, Congzheng Song, Emiliano De Cristofaro, and Vitaly Shmatikov. 2018. Inference attacks against collaborative learning. *CoRR* abs/1805.04049 (2018). arxiv:1805.04049 <http://arxiv.org/abs/1805.04049>.
- [44] Payman Mohassel and Peter Rindal. 2018. ABY3: A mixed protocol framework for machine learning. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security (CCS’18)*. ACM, New York, NY, 35–52. DOI : <https://doi.org/10.1145/3243734.3243760>
- [45] Payman Mohassel, Mike Rosulek, and Ye Zhang. 2015. Fast and secure three-party computation: The garbled circuit approach. In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security (CCS’15)*. ACM, New York, NY, 591–602. DOI : <https://doi.org/10.1145/2810103.2813705>
- [46] Payman Mohassel and Yupeng Zhang. 2017. SecureML: A system for scalable privacy-preserving machine learning. In *IEEE Symposium on Security and Privacy*. IEEE Computer Society, 19–38.
- [47] Payman Mohassel and Yupeng Zhang. 2017. SecureML: A system for scalable privacy-preserving machine learning. *IACR Cryptology ePrint Archive* 2017 (2017), 396.
- [48] Valeria Nikolaenko, Udi Weinsberg, Stratis Ioannidis, Marc Joye, Dan Boneh, and Nina Taft. 2013. Privacy-preserving ridge regression on hundreds of millions of records. In *Proceedings of the 2013 IEEE Symposium on Security and Privacy (SP’13)*. IEEE Computer Society, Washington, DC, 334–348. DOI : <https://doi.org/10.1109/SP.2013.30>
- [49] Richard Nock, Stephen Hardy, Wilko Henecka, Hamish Ivey-Law, Giorgio Patrini, Guillaume Smith, and Brian Thorne. 2018. Entity resolution and federated learning get a federated resolution. *CoRR* abs/1803.04035 (2018). arxiv:1803.04035 <http://arxiv.org/abs/1803.04035>.
- [50] Simo Jialin Pan and Qiang Yang. 2010. A survey on transfer learning. *IEEE Trans. Knowl. Data Eng.* 22, 10 (Oct. 2010), 1345–1359. DOI : <https://doi.org/10.1109/TKDE.2009.191>
- [51] Le Trieu Phong, Yoshinori Aono, Takuya Hayashi, Lihua Wang, and Shiho Moriai. 2018. Privacy-preserving deep learning via additively homomorphic encryption. *IEEE Trans. Information Forensics and Security* 13, 5 (2018), 1333–1345.
- [52] M. Sadegh Riazi, Christian Weinert, Oleksandr Tkachenko, Ebrahim M. Songhori, Thomas Schneider, and Farinaz Koushanfar. 2018. Chameleon: A hybrid secure computation framework for machine learning applications. *CoRR* abs/1801.03239 (2018).

- [53] R. L. Rivest, L. Adleman, and M. L. Dertouzos. 1978. On data banks and privacy homomorphisms. *Foundations of Secure Computation* 4, 11 (1978), 169–179.
- [54] Bita Darvish Rouhani, M. Sadegh Riazi, and Farinaz Koushanfar. 2017. DeepSecure: Scalable provably-secure deep learning. *CoRR* abs/1705.08963 (2017). arxiv:1705.08963 <http://arxiv.org/abs/1705.08963>.
- [55] Ashish P. Sanil, Alan F. Karr, Xiaodong Lin, and Jerome P. Reiter. 2004. Privacy preserving regression modelling via distributed computation. In *Proceedings of the 10th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'04)*. ACM, New York, NY, 677–682. DOI: <https://doi.org/10.1145/1014052.1014139>
- [56] Monica Scannapieco, Ilya Figotin, Elisa Bertino, and Ahmed K. Elmagarmid. 2007. Privacy preserving schema and data matching. In *Proceedings of the 2007 ACM SIGMOD International Conference on Management of Data (SIGMOD'07)*. ACM, New York, NY, 653–664. DOI: <https://doi.org/10.1145/1247480.1247553>
- [57] Amit P. Sheth and James A. Larson. 1990. Federated database systems for managing distributed, heterogeneous, and autonomous databases. *ACM Comput. Surv.* 22, 3 (Sept. 1990), 183–236. DOI: <https://doi.org/10.1145/96602.96604>
- [58] Reza Shokri and Vitaly Shmatikov. 2015. Privacy-preserving deep learning. In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security (CCS'15)*. ACM, New York, NY, 1310–1321. DOI: <https://doi.org/10.1145/2810103.2813687>
- [59] David Silver, Aja Huang, Christopher J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. 2016. Mastering the game of Go with deep neural networks and tree search. *Nature* 529 (2016), 484–503. <http://www.nature.com/nature/journal/v529/n7587/full/nature16961.html>.
- [60] Virginia Smith, Chao-Kai Chiang, Maziar Sanjabi, and Ameet S. Talwalkar. 2017. Federated multi-task learning. In *Advances in Neural Information Processing Systems 30*, I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett (Eds.). Curran Associates, Inc., 4424–4434. <http://papers.nips.cc/paper/7029-federated-multi-task-learning.pdf>.
- [61] Shuang Song, Kamalika Chaudhuri, and Anand D. Sarwate. 2013. Stochastic gradient descent with differentially private updates. *2013 IEEE Global Conference on Signal and Information Processing* (2013). IEEE, 245–248.
- [62] Lili Su and Jiaming Xu. 2018. Securing distributed machine learning in high dimensions. *CoRR* abs/1804.10140 (2018). arxiv:1804.10140 <http://arxiv.org/abs/1804.10140>.
- [63] Latanya Sweeney. 2002. K-anonymity: A model for protecting privacy. *Int. J. Uncertain. Fuzziness Knowl.-Based Syst.* 10, 5 (Oct. 2002), 557–570. DOI: <https://doi.org/10.1142/S0218488502001648>
- [64] Jaideep Vaidya and Chris Clifton. 2004. Privacy preserving naïve Bayes classifier for vertically partitioned data. In *Proceedings of the 4th SIAM Conference on Data Mining, 2004*. 330–334.
- [65] Jaideep Vaidya and Chris Clifton. 2002. Privacy preserving association rule mining in vertically partitioned data. In *Proceedings of the 8th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'02)*. ACM, New York, NY, 639–644. DOI: <https://doi.org/10.1145/775047.775142>
- [66] Jaideep Vaidya and Chris Clifton. 2003. Privacy-preserving K-means clustering over vertically partitioned data. In *Proceedings of the 9th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'03)*. ACM, New York, NY, 206–215. DOI: <https://doi.org/10.1145/956750.956776>
- [67] Jaideep Vaidya and Chris Clifton. 2005. Privacy-preserving decision trees over vertically partitioned data. In *Data and Applications Security XIX*, Sushil Jajodia and Duminda Wijesekera (Eds.). Springer, Berlin, 139–152.
- [68] Li Wan, Wee Keong Ng, Shuguo Han, and Vincent C. S. Lee. 2007. Privacy-preservation for gradient descent methods. In *Proceedings of the 13th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'07)*. ACM, New York, NY, 775–783. DOI: <https://doi.org/10.1145/1281192.1281275>
- [69] Shiqiang Wang, Tiffany Tuor, Theodoros Salonidis, Kin K. Leung, Christian Makaya, Ting He, and Kevin Chan. 2018. When edge meets learning: Adaptive control for resource-constrained distributed machine learning. *CoRR* abs/1804.05271 (2018). arxiv:1804.05271 <http://arxiv.org/abs/1804.05271>.
- [70] Wikipedia. 2018. Facebook–Cambridge Analytica Data Scandal. [https://en.wikipedia.org/wiki/Facebook–Cambridge\\_Analytica%\\_data\\_scandal](https://en.wikipedia.org/wiki/Facebook–Cambridge_Analytica%_data_scandal).
- [71] Qiang Yang, Yang Liu, Tianjian Chen, and Yongxin Tong. 2018. Federated learning. *Communications of the CCF* 14, 11 (2018), 49–55.
- [72] Andrew C. Yao. 1982. Protocols for secure computations. In *Proceedings of the 23rd Annual Symposium on Foundations of Computer Science (SFCS'82)*. IEEE Computer Society, Washington, DC, 160–164. <http://dl.acm.org/citation.cfm?id=1382436.1382751>.
- [73] Hwanjo Yu, Xiaoqian Jiang, and Jaideep Vaidya. 2006. Privacy-preserving SVM using nonlinear kernels on horizontally partitioned data. In *Proceedings of the 2006 ACM Symposium on Applied Computing (SAC'06)*. ACM, New York, NY, 603–610. DOI: <https://doi.org/10.1145/1141277.1141415>

- [74] Hwanjo Yu, Jaideep Vaidya, and Xiaoqian Jiang. 2006. Privacy-preserving SVM classification on vertically partitioned data. In *Proceedings of the 10th Pacific-Asia Conference on Advances in Knowledge Discovery and Data Mining (PAKDD'06)*. Springer, Berlin, 647–656. DOI:[https://doi.org/10.1007/11731139\\_74](https://doi.org/10.1007/11731139_74)
- [75] Jiawei Yuan and Shucheng Yu. 2014. Privacy preserving back-propagation neural network learning made practical with cloud computing. *IEEE Trans. Parallel Distrib. Syst.* 25, 1 (Jan. 2014), 212–221. DOI:<https://doi.org/10.1109/TPDS.2013.18>
- [76] Qingchen Zhang, Laurence T. Yang, and Zhikui Chen. 2016. Privacy preserving deep computation model on cloud for big data feature learning. *IEEE Trans. Comput.* 65, 5 (May 2016), 1351–1362. DOI:<https://doi.org/10.1109/TC.2015.2470255>
- [77] Yue Zhao, Meng Li, Liangzhen Lai, Naveen Suda, Damon Civin, and Vikas Chandra. 2018. Federated Learning with Non-IID Data. arxiv:cs.LG/1806.00582

Received September 2018; accepted November 2018